



Apigee Micro Gateways on Kubernetes

Build once , run anywhere



Updated Date: 2/15/2018



Agenda

Day 1

Introduction

- ✓ About the Pilot
- ✓ Outcome

Kubernetes

- ✓ Architecture
- ✓ Basic Commands
- ✓ PODS
- ✓ Deployment
- ✓ Services
- ✓ Primitives
- ✓ Networking
- ✓ Ingress
- ✓ Name Spaces
- Configuration
- ✓ Service Discovery

Hands On

- ✓ Kubernetes

Day 2

Edge Micro Gateway

- Architecture
- Deployment Options
- Gateway as an Ingress
- Security
- Analytics

CICD

- Jenkins Jobs
 - Build
 - Multi Cluster Deploy

Pilot Knowledge Transfer

- Micro Gateway Container
- Micro Service
- KOPS on AWS
- Kubernetes on Azure

Hands On

- Edge Micro Gateway



Pilot Details

Scope :

- Create 3 node Kubernetes cluster on AWS. Cluster should have 1 master and 2 worker nodes
- Create 3 node Kubernetes cluster on Azure. Cluster should have 1 master and 2 worker nodes
- Create and Configure ingress controllers on both AWS and Azure
- Create a total of 4 different work spaces
- Build and configure Apigee micro gateway (AMG) container capable of being initialized with environment specific information
- Create and deploy service composed of AMG and helloworld microservices as shown in the architecture to the right in parallel on AWS and Azure
- Demonstrate that the architecture is scalable, portable and cloud agnostic

Deliverables:

1. Kubernetes (k8s) Cluster on AWS with 1 master and 2 worker nodes with ingress controllers
2. Kubernetes Cluster on Azure with 1 master and 2 worker nodes with ingress controllers
3. Jenkins Server with Two (2) Jenkins Jobs to deploy Service to AWS and Azure k8s
4. Git repo to build a AMG image capable of taking environment information as a runtime variable
5. Demo showing the following
 1. Solution is cloud agnostic
 2. Solution is scalable
6. Documentation of Steps to repeat items #1, #2 and #3



Outcome

- Deploy a Micro Gateway on AWS and AZURE using common manifest
- Deploy Micro Service on AWS and AZURE using common manifest
- Observe Scaling and Load balancing on both clusters
- Create 1 proxy for both Micro gateways with a cloud and cluster agnostic target
- Observe Cloud agnostic Routing
- Provide Knowledge Transfer



Day 1





Part 1

- Introduction
- What is it Really?
- Installation



Introduction

- Kubernetes is an open-source software for automating deployment, scaling, and management of containerized workloads.
- Builds on 15 years of experience at Google.



Alternatives (not an exhaustive list)

- Docker Swarm
- Cloud Foundry
- Apache Mesos
- Rancher
- Hashicorp Nomad



Why are we talking about it today?

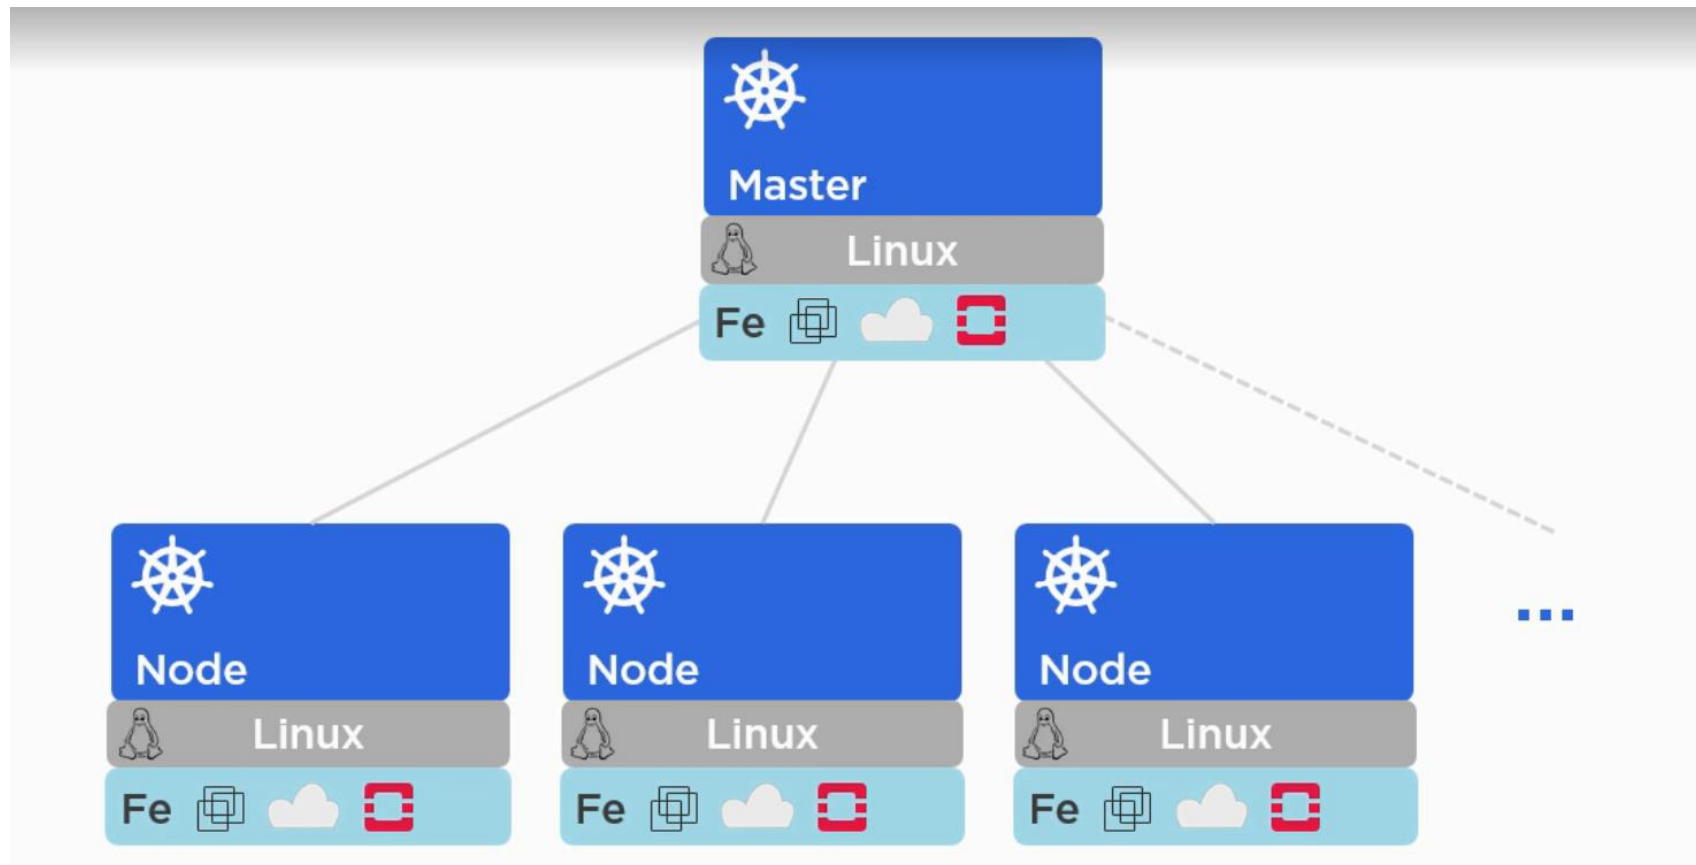
Because its your target environment where a majority of your apps and integrations will run and it runs everywhere...

- Amazon Cloud
- Azure Cloud
- Runs on VMWare

Finally fulfills the promise of “Code Once, Run Anywhere..”



Why Kubernetes, what is Kubernetes?




What :
Cluster Manager
Container Scheduler

Why ?
Geared towards
scale
Lots of HA features
Deployed on Bare
Metal or VM's



Installation Options

- Minikube (Easiest to set up)



- KOPS (AWS)
- ACS (Azure)

Used for Pilot

- kubeadm
- hyperkube
- Bootkube



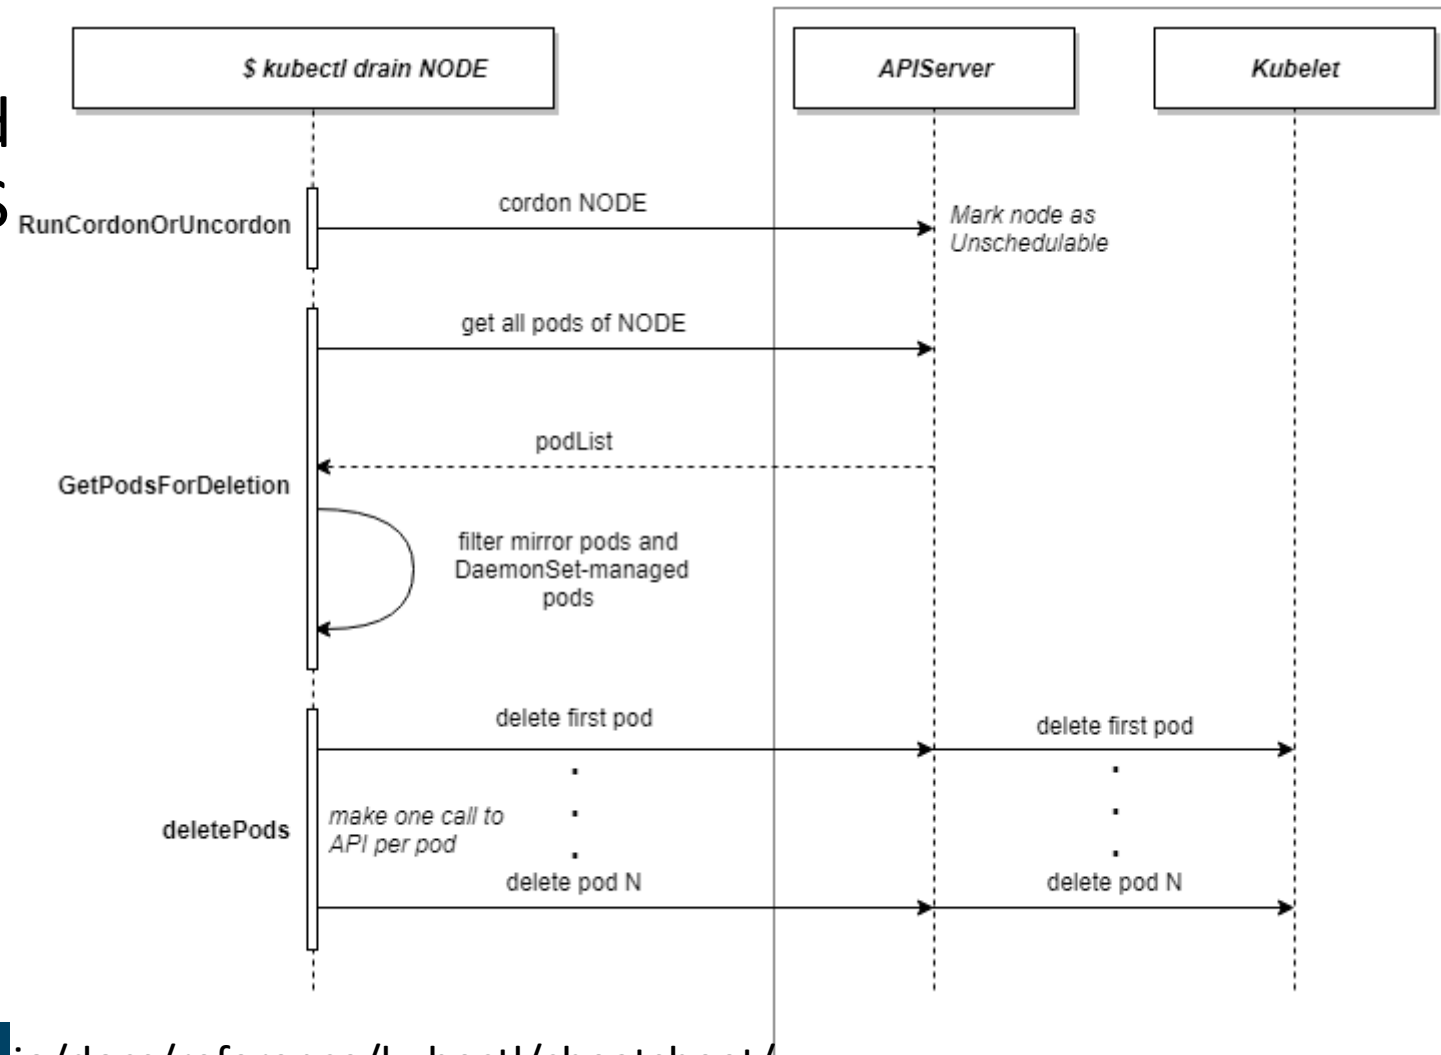
Part 2

- Kubectl
 - The life of an API
- Kubeconfig
- Kubernetes Primitives



Role of Kubectl

- Command-line tool configured to communicate with your k8S cluster
- kubectl
 - Basic commands
 - Exec [For debugging]
- How does kubectl work?
 - Ex : Drain a NODE





Life of an API request

- The request is received on a TLS secured API Server
- The API Server
 - Authenticates the request using one of
 - Certificates
 - Tokens
 - Basic Auth
 - Open Id
 - Web Hooks
 - Authorizes the request (RBAC based on group memberships)
 - Runs admission control policies
- The actual work is done



Role of Kubeconfig

- The file used to configure access to a K8S cluster
- Provides information about
 - Cluster Location
 - Users
 - Context - A combination of cluster + user [+ namespace]

```
apiVersion: v1
kind: Config
clusters:
- cluster:
  certificate-authority: ssl/ca.pem
  server: https://172.17.4.101:443
  name: vagrant-multi-cluster
contexts:
- context:
  cluster: vagrant-multi-cluster
  namespace: default
  user: vagrant-multi-admin
  name: vagrant-multi
users:
- name: vagrant-multi-admin
  user:
  client-certificate: ssl/admin.pem
  client-key: ssl/admin-key.pem
current-context: vagrant-multi
```

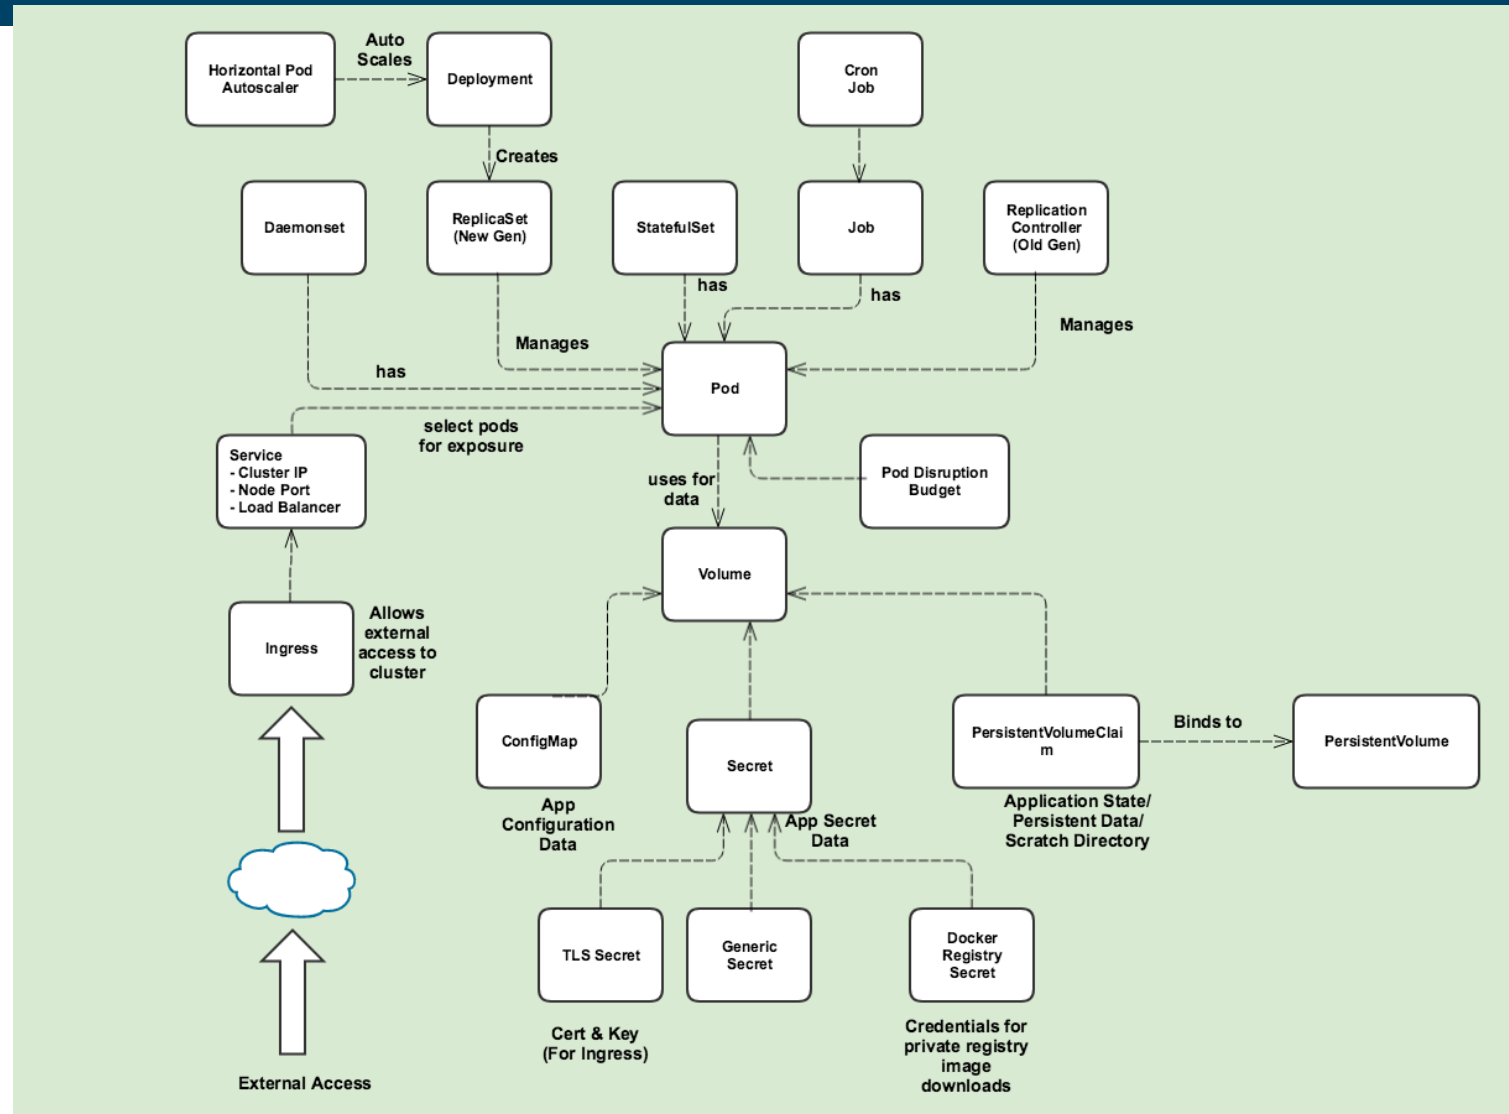


Kubernetes Primitives – Developers Perspective

Pod Autoscaler
Deployment
Daemonset
ReplicaSet
Stateful Set
Job
Cron Job
Replication Controller

Pod
Pod Disruption Budget
Volume
ConfigMap
Secret (* Secret)
Persistent Volume Claim

Ingress
Service
- Cluster IP
- Node Port
- Load Balancer





What's is a Pod ?



VM



Container

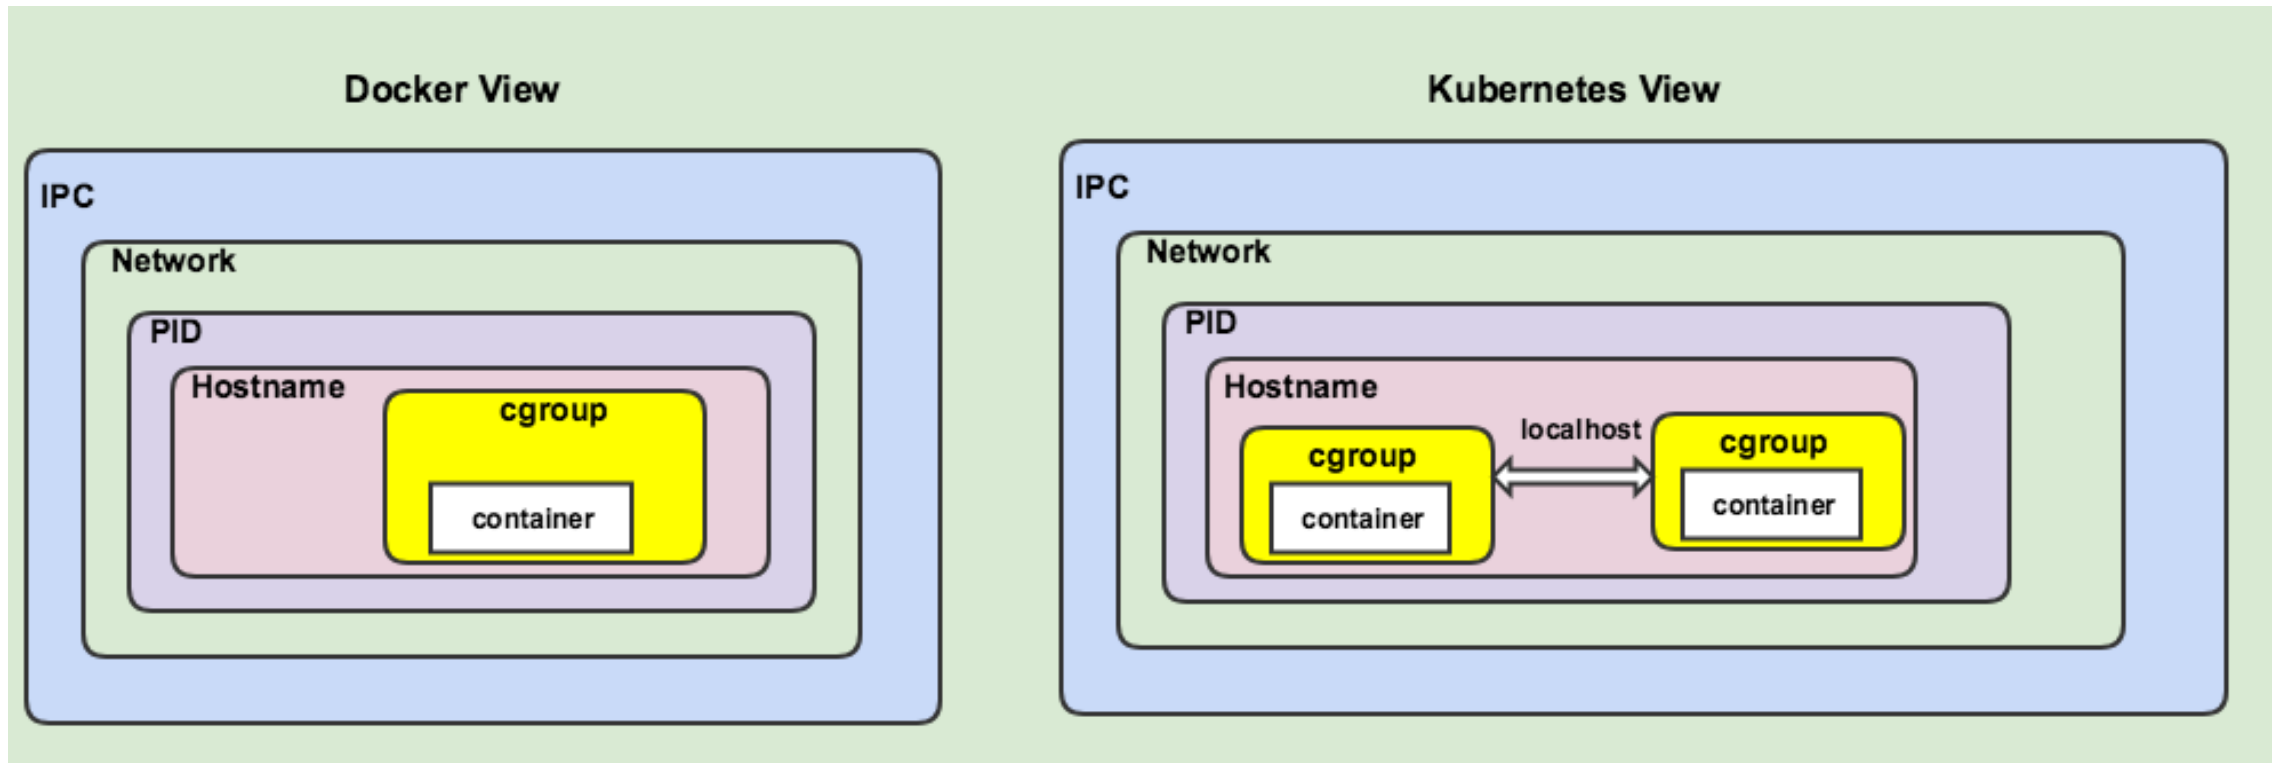


Pod

Atomic units of scheduling



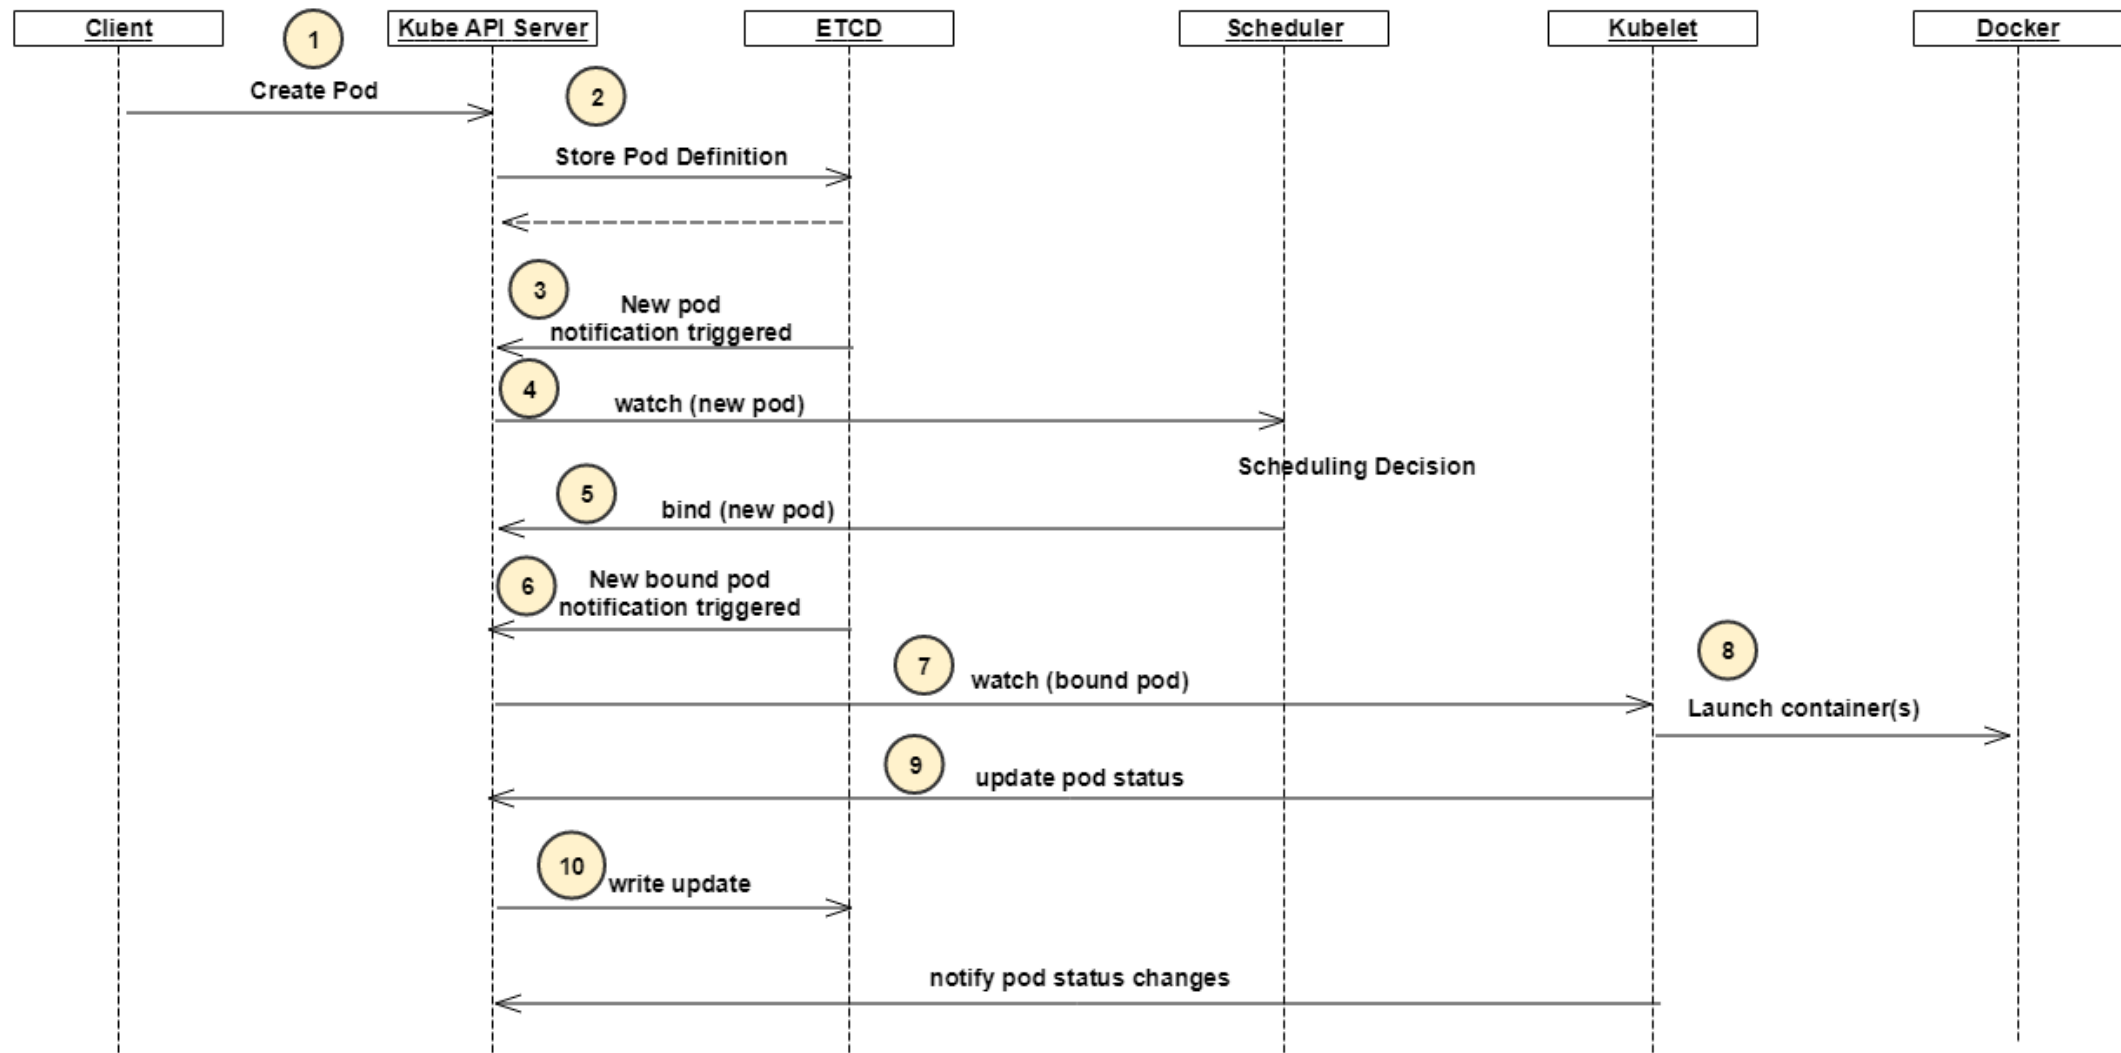
What's in a Pod ?



cgroups (abbreviated from control groups) is a Linux kernel feature that limits, accounts for, and isolates the resource usage (CPU, memory, disk I/O, network, etc.) of a collection of processes



Scheduling a Pod





Kubernetes – Key takeaways (Recap)

- Actual State vs Desired State
- Based on the concept of Resources
 - Each Resource is defined as an object in a YAML file
 - A Controller inspects each resource and manages its lifecycle
 - Roles and permissions are managed around CRUD permissions on resources
- Imperative vs Declarative commands



- 15 Minute Break



Part 3

- Namespaces
- Replica Sets
- Deployments
- Labels and Services
- Accessing applications
 - Via services
 - Via Ingress



Namespaces

Namespaces are intended to isolate multiple groups / teams and give the access to a set of resources. Each namespace can have quotas

Think: Virtual Clusters

Note:

- Namespaces are flat and cannot be nested
- Namespaces DO NOT provide network isolation

Hands On:

- Creating a namespace
- Using a namespace



Replica Sets (Replication Controller)

A replica set ensures that a specified number of pod “replicas” are running at any one time.

In other words, A replica set makes sure that a pod or homogeneous set of pods are always up and available.

Hands On:
Replica Sets



Deployments

- Allow Server side update of pods at a specified rate
- Generate Replica Sets
- Allow for Deployment Patterns
 - Blue Green
 - Canary

Hands On:

- Scale up and Down
- Blue Green
- Canary
- Canary (Dedicated Access)



Labels and Services

- Labels select objects
- Services expose applications

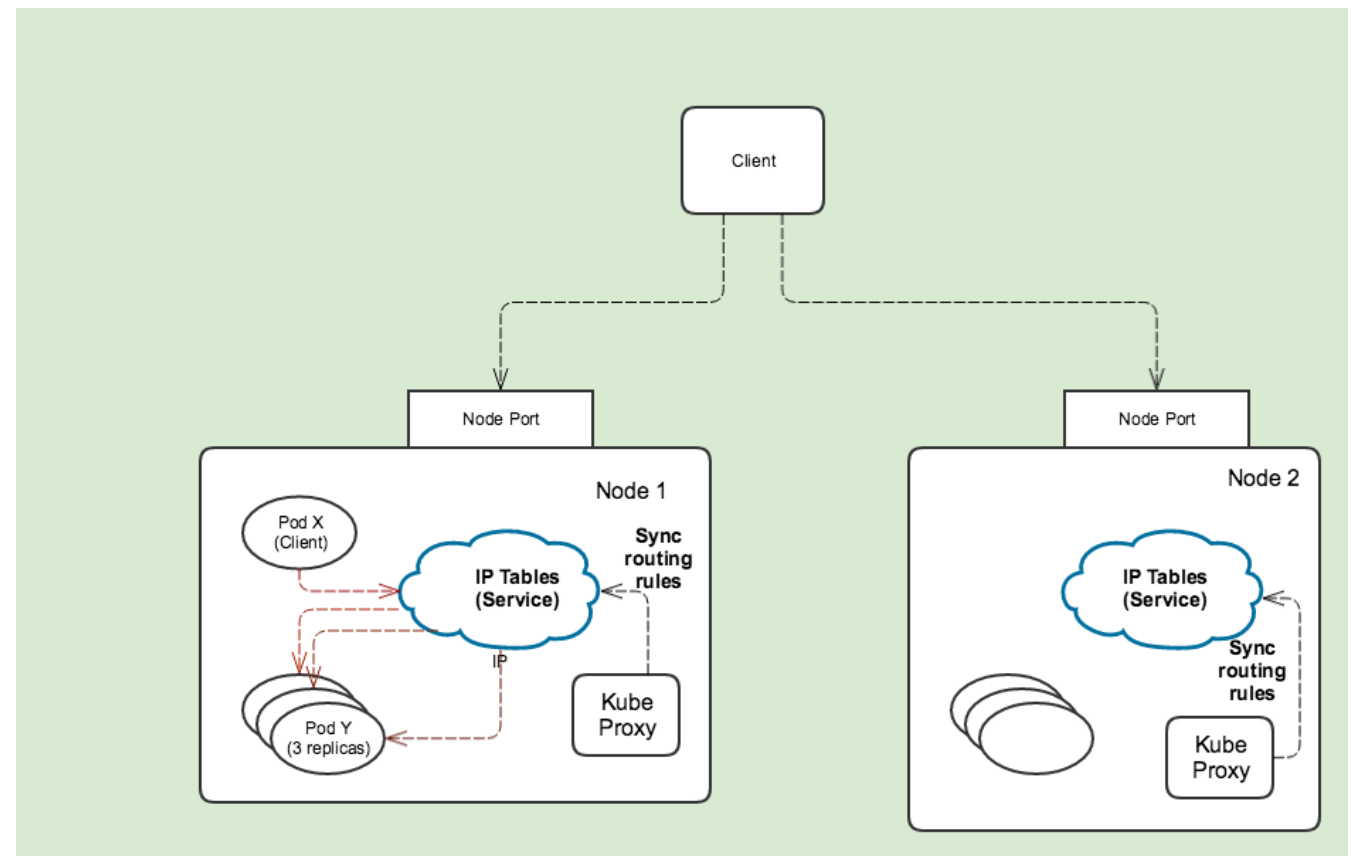
Hands On:

- Labels
- Services
- Load Balancing



Services – Accessing your applications

- Provide a stable virtual endpoint to ephemeral pods
- Not tied to any network interface
- Based purely on IP Tables (as of v 1.9)





Service Types

- **Cluster IP** – Default, Provides internal access to the endpoint (unless manually created to point to an external endpoint)
- **Node Port** – Exposes the same high range port for every host
- **Load Balancer** – Native solution to expose services on cloud providers



Scaling and Rolling updates - Deployments

- Demos
 - Scale up
 - Guarding against bad deployments
 - Rolling back deployments



Ingress

- A reverse proxy like functionality giving you access to your applications
- Collection of rules that allow inbound access to the cluster services

Think : Nginx (But less configuration and maintenance)



Anatomy of an Ingress

- Ingress Resources (Defined by Developers per access point)
- Ingress Controller (Created by Cluster/Namespace Administrator once)

Demos:

- Ingress Controller
- Ingress Resources (Vanilla)
- Ingress Resources with TLS Termination



Part 4

- Volumes
 - Secrets
 - Config Maps
 - Persistent Volume and Persistent Volume Claim



Config Maps

- Stores non sensitive configuration data
- Demo



Secrets

- Store Secret Data
 - Generic Secret
 - TLS Secret
 - Docker Registry Secret



Volumes

- Persistent Volumes (Storage Abstraction)
- Persistent Volume Claim (Pod's claim to a volume)



- 15 Minute Break

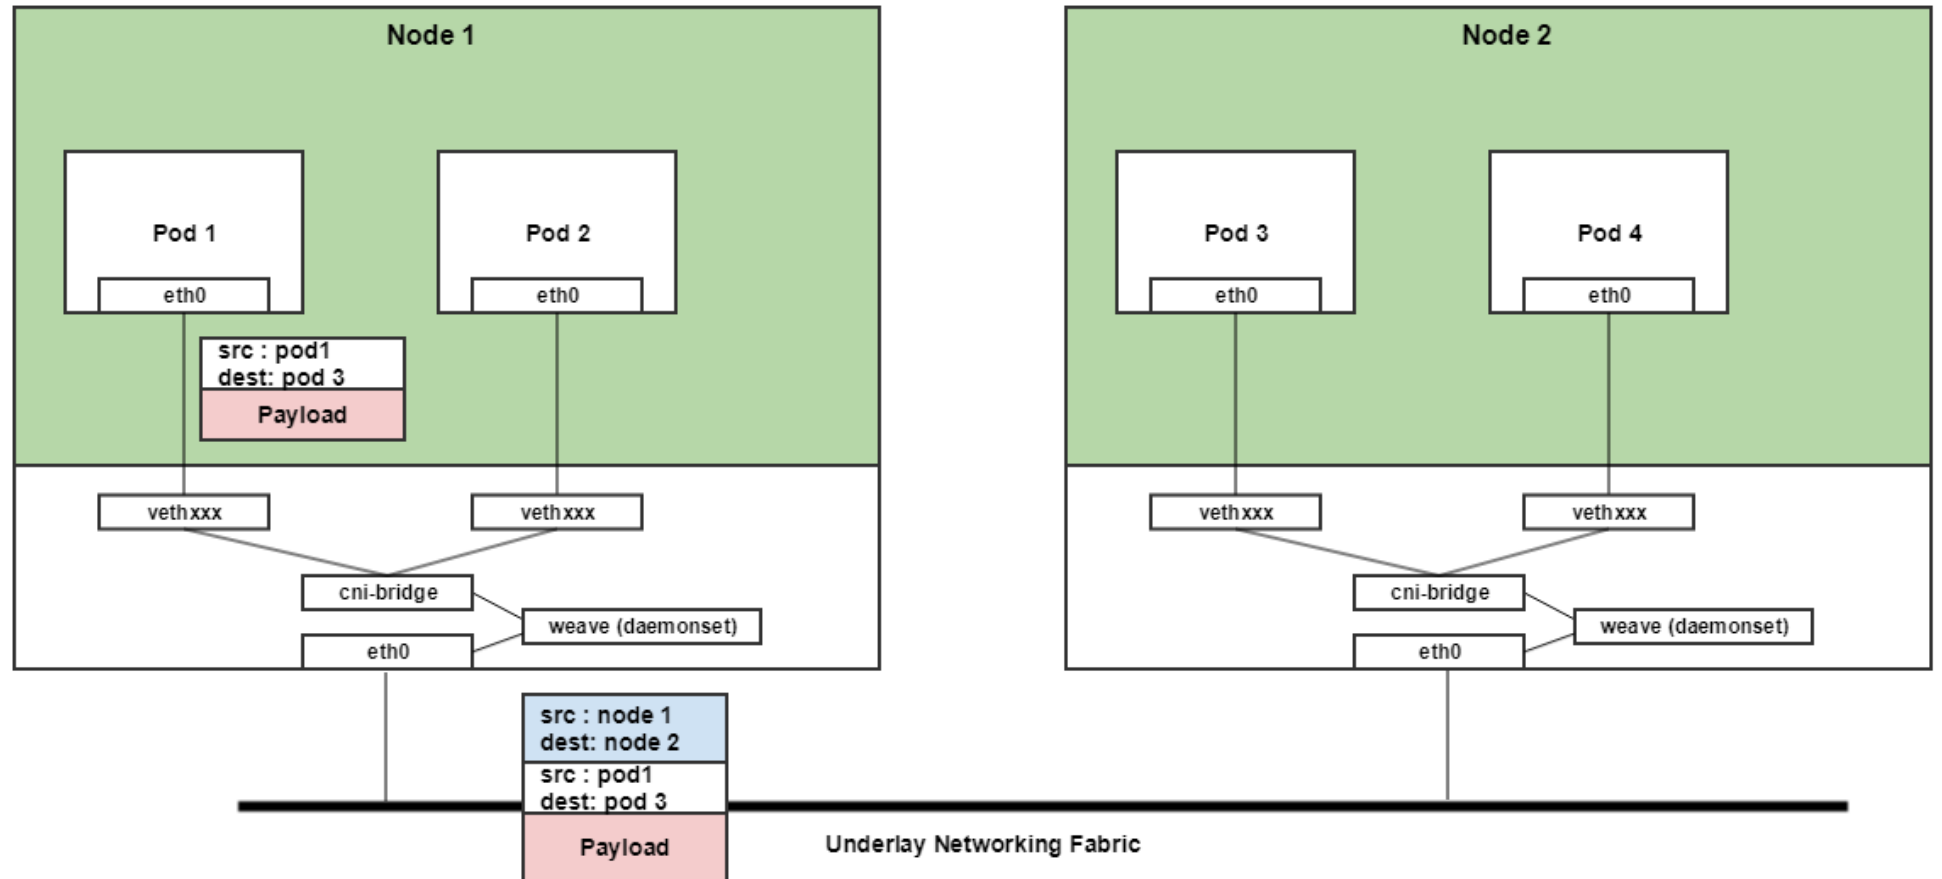


Part 5

- Networking Primer



Kubernetes Networking Primer





Thank you